

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import warnings
        5 warnings.filterwarnings("ignore")
```

```
In [5]: 1 data=pd.read_csv(r"C:\Users\Admin\Downloads\diabetes.csv")
```

```
In [6]: 1 data
```

Out[6]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | 0.340 | 27 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | 0.349 | 47 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | 0.315 | 23 |

768 rows × 9 columns



```
In [7]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                 768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                    768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                    768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
]: 1
```

```
]:
```

r i s t

```
In [8]: data.describe()
```

```
Out[8]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigree |
|-------|-------------|------------|---------------|---------------|------------|------------|------------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | 0.398664 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | 0.327144 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.078000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | 0.171000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | 0.245000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | 0.349000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | 0.672000 |

```
In [9]: 1 data.head()
```

```
Out[9]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |

```
In [10]: 1 data.tail()
```

```
Out[10]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | 0.340 | 27 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | 0.349 | 47 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | 0.315 | 23 |

```
In [11]: 1 data["Outcome"].value_counts(normalize=True)
```

```
Out[11]: Outcome
0    0.651042
1    0.348958
Name: proportion, dtype: float64
```

```
In [12]: 1 x=data.drop(["Outcome"],axis=1)
```

```
In [13]: 1 y=data["Outcome"]
```

In []:

1

]:

r i s t

14 x

Out[14]

| | Pregnancies | Glucose | BloodPressu e | SkinTh ckness | In ulin | BMI | DiabetesPedigreeFunc ion | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | 0.340 | 27 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | 0.349 | 47 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | 0.315 | 23 |

768 rows × 8 columns

In [15]:

1 y

Out[15]:

| | |
|-----|---|
| 0 | 1 |
| 1 | 0 |
| 2 | 1 |
| 3 | 0 |
| 4 | 1 |
| ... | |
| 763 | 0 |
| 764 | 0 |
| 765 | 0 |
| 766 | 1 |
| 767 | 0 |

Name: Outcome, Length: 768, dtype: int64

In [16]:

1 from sklearn.model_selection import train_test_split

In [17]:

1 train_x, test_x, train_y , test_y =train_test_split(x,y,random_state =56)

In []:

1

]:

r i s t

18

train_x

Out[18]

| | Pregnancies | Glucose | BloodPressu e | SkinTh ckness | In ulin | BMI | DiabetesPedigreeFunc ion | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 536 | 0 | 105 | 90 | 0 | 0 | 29.6 | 0.197 | 46 |
| 547 | 4 | 131 | 68 | 21 | 166 | 33.1 | 0.160 | 28 |
| 307 | 0 | 137 | 68 | 14 | 148 | 24.8 | 0.143 | 21 |
| 45 | 0 | 180 | 66 | 39 | 0 | 42.0 | 1.893 | 25 |
| 196 | 1 | 105 | 58 | 0 | 0 | 24.3 | 0.187 | 21 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 235 | 4 | 171 | 72 | 0 | 0 | 43.6 | 0.479 | 26 |
| 418 | 1 | 83 | 68 | 0 | 0 | 18.2 | 0.624 | 27 |
| 192 | 7 | 159 | 66 | 0 | 0 | 30.4 | 0.383 | 36 |
| 399 | 3 | 193 | 70 | 31 | 0 | 34.9 | 0.241 | 25 |
| 484 | 0 | 145 | 0 | 0 | 0 | 44.2 | 0.630 | 31 |

576 rows × 8 columns



In [19]:

1 test_x

Out[19]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 123 | 5 | 132 | 80 | 0 | 0 | 26.8 | 0.186 | 69 |
| 295 | 6 | 151 | 62 | 31 | 120 | 35.5 | 0.692 | 28 |
| 370 | 3 | 173 | 82 | 48 | 465 | 38.4 | 2.137 | 25 |
| 300 | 0 | 167 | 0 | 0 | 0 | 32.3 | 0.839 | 30 |
| 155 | 7 | 152 | 88 | 44 | 0 | 50.0 | 0.337 | 36 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 443 | 8 | 108 | 70 | 0 | 0 | 30.5 | 0.955 | 33 |
| 134 | 2 | 96 | 68 | 13 | 49 | 21.1 | 0.647 | 26 |
| 181 | 0 | 119 | 64 | 18 | 92 | 34.9 | 0.725 | 23 |
| 588 | 3 | 176 | 86 | 27 | 156 | 33.3 | 1.154 | 52 |
| 737 | 8 | 65 | 72 | 23 | 0 | 32.0 | 0.600 | 42 |

192 rows × 8 columns



In [20]:

1 from sklearn.preprocessing import MinMaxScaler

In [23]:

1 scaler=MinMaxScaler()
2 scaler

Out[23]:

| |
|----------------|
| ▼ MinMaxScaler |
| MinMaxScaler() |

In []: 1

```
24 cols=train_x.columns
2 cols
```

Out[24]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'], dtype='object')

```
In [25]: 1 train_x_scaled=scaler.fit_transform(train_x)
2 train_x_scaled
```

Out[25]: array([[0. , 0.53030303, 0.73770492, ..., 0.44113264, 0.05081127, 0.41666667],
[0.23529412, 0.66161616, 0.55737705, ..., 0.49329359, 0.03501281, 0.11666667],
[0. , 0.69191919, 0.55737705, ..., 0.36959762, 0.02775406, 0.],
...,
[0.41176471, 0.8030303 , 0.54098361, ..., 0.45305514, 0.13023057, 0.25],
[0.17647059, 0.97474747, 0.57377049, ..., 0.52011923, 0.06959863, 0.06666667],
[0. , 0.73232323, 0. , ..., 0.65871833, 0.23569599, 0.16666667]])

```
In [26]: 1 train_x_scaled=pd.DataFrame(train_x_scaled,columns=cols)
```

```
In [27]: 1 train_x_scaled
```

Out[27]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction |
|-----|-------------|----------|---------------|---------------|----------|----------|--------------------------|
| 0 | 0.000000 | 0.530303 | 0.737705 | 0.000000 | 0.000000 | 0.441133 | 0.05081 |
| 1 | 0.235294 | 0.661616 | 0.557377 | 0.212121 | 0.196217 | 0.493294 | 0.03501 |
| 2 | 0.000000 | 0.691919 | 0.557377 | 0.141414 | 0.174941 | 0.369598 | 0.02775 |
| 3 | 0.000000 | 0.909091 | 0.540984 | 0.393939 | 0.000000 | 0.625931 | 0.77497 |
| 4 | 0.058824 | 0.530303 | 0.475410 | 0.000000 | 0.000000 | 0.362146 | 0.04654 |
| ... | ... | ... | ... | ... | ... | ... | . |
| 571 | 0.235294 | 0.863636 | 0.590164 | 0.000000 | 0.000000 | 0.649776 | 0.17122 |
| 572 | 0.058824 | 0.419192 | 0.557377 | 0.000000 | 0.000000 | 0.271237 | 0.23313 |
| 573 | 0.411765 | 0.803030 | 0.540984 | 0.000000 | 0.000000 | 0.453055 | 0.13023 |
| 574 | 0.176471 | 0.974747 | 0.573770 | 0.313131 | 0.000000 | 0.520119 | 0.06959 |
| 575 | 0.000000 | 0.732323 | 0.000000 | 0.000000 | 0.000000 | 0.658718 | 0.23569 |

576 rows × 8 columns



```
In [28]: 1 from sklearn.linear_model import LogisticRegression as LogReg
```

```
In [29]: 1 logreg=LogReg()
```

```
In [30]: 1 logreg.fit(train_x,train_y)
```

Out[30]:

LogisticRegression

LogisticRegression()

In []: 1

```
31 train_predict=logreg.predict(train_x)
2 test_predict=logreg.predict(test_x)
```

In [32]: 1 train_predict

```
Out[32]: array([0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0,
1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0,
0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0,
1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0,
0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0,
0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0,
0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0,
1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
0, 1, 1, 1], dtype=int64)
```

In [33]: 1 test_predict

```
Out[33]: array([0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0,
0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0,
0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0], dtype=int64)
```

In [34]: 1 from sklearn.metrics import f1_score, confusion_matrix, roc_auc_score, roc_curve

In [36]: 1 f1_score(train_predict,train_y)

Out[36]: 0.6304347826086956

In [37]: 1 f1_score(test_predict,test_y)

Out[37]: 0.7008547008547009

In [38]: 1 conf1=confusion_matrix(train_y,train_predict)

In [39]: 1 conf1

```
Out[39]: array([[324, 42],
[115, 115]], dtype=int64)
```

In []: 1

```
40 conf=confusion_matrix(test_y,test_predict)
```

In [41]: 1 conf

Out[41]: array([[116, 18],
[17, 41]], dtype=int64)

```
In [42]: 1 true_negative =conf[0][0]  
2 false_negative =conf[1][0]  
3 false_positive =conf[0][1]  
4 true_positive =conf[1][1]
```

```
In [43]: 1 Accuracy = (true_positive + true_negative) / (true_positive +false_positive + f  
2 Accuracy  
3  
4 # Precision  
5 Precision = true_positive/(true_positive+false_positive)  
6 Precision  
7  
8 # Recall  
9 Recall = true_positive/(true_positive+false_negative)  
10 Recall  
11  
12 # F1 Score  
13 F1_Score = 2*(Recall * Precision) / (Recall + Precision)  
14 F1_Score
```

Out[43]: 0.7008547008547009

In [44]: 1 Accuracy

Out[44]: 0.8177083333333334

In [45]: 1 Precision

Out[45]: 0.6949152542372882

In [46]: 1 Recall

Out[46]: 0.7068965517241379

In [47]: 1 F1_Score

Out[47]: 0.7008547008547009

```
In [48]: 1 auc_score=roc_auc_score(test_y,test_predict)
```

```
In [49]: 1 fpr,tpr,thresholds=roc_curve(test_y,test_predict)
```


```
In [50]: 1 thresholds
```

Out[50]: array([inf, 1., 0.])

In []: 1

51

```
plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (area = 0.79)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



In []:

1

1